

Implementation of A*

from Red Blob Games

*Jul 2014, then Feb 2016, Nov 2018,
Oct 2020*

Table of Contents

This article is a companion guide to my [introduction to A*](#), where I explain how the algorithms work. On this page I show how to implement Breadth-First Search, Dijkstra's Algorithm, Greedy Best-First Search, and A*. I try to keep the code here simple.

Graph search is a family of related algorithms. There are *lots* of variants of the algorithms, and lots of variants in implementation. Treat the code on this page as a starting point, not as a final version of the algorithm that works for all situations.

1

- 1. Python Implementation
 - 1.1. Breadth First Search
 - 1.2. Early Exit
 - 1.3. Dijkstra's Algorithm
 - 1.4. A* Search
- 2. C++ Implementation
 - 2.1. Breadth First Search
 - 2.2. Early Exit
 - 2.3. Dijkstra's Algorithm
 - 2.4. A* Search
 - 2.5. Production code
- 3. C# Implementation
- 4. Algorithm changes
- 5. Optimizations
 - 5.1. Graph
 - 5.2. Queue
 - 5.3. Search
 - 5.4. Integer locations
- 6. Troubleshooting
 - 6.1. Wrong paths
 - 6.2. Ugly paths
- 7. Vocabulary
- 8. More reading

Python Implementation

#

I explain most of the code below. There are a few extra bits that you can find in [implementation.py](#). These use **Python 3** so if you use Python 2, you will need to remove type annotations, change the `super()` call, and change the `print` function to work with Python 2.

1.1 Breadth First Search

#

Let's implement Breadth First Search in Python. The main article shows the Python code for the search algorithm, but we also need to define the graph it works on. These are the abstractions I'll use:

Graph

a data structure that can tell me the `neighbors` for each graph location (see [this tutorial](#)). A *weighted* graph also gives a `cost` of moving along an edge.

Locations

a simple value (int, string, tuple, etc.) that *labels* locations in the graph. These are not necessarily locations on the map. They may include additional information such as direction, fuel, lane, or inventory, depending on the problem being solved.

Search

an algorithm that takes a graph, a starting graph location, and optionally a goal graph location, and calculates some useful information (reached, parent pointer, distance) for some or all graph locations.

Queue

a data structure used by the search algorithm to decide the order in which to process the graph locations.

In the main article, I focused on **search**. On this page, I'll fill in the rest of the details to make complete working programs. Let's start with a **graph**. What does a graph look like? It's a **location** type along with a class with a method to get neighboring locations:

```
Location = TypeVar('Location')
class Graph(Protocol):
```

```
def neighbors(self, id: Location) -> List[Location]: pass
```


`Graph` is the interface that the search algorithms will want. Here's an implementation go to with it:

```
class SimpleGraph:
    def __init__(self):
        self.edges: Dict[Location, List[Location]] = {}


    def neighbors(self, id: Location) -> List[Location]:
        return self.edges[id]
```

Yep, that's all we need! You may be asking, where's the `Node` object? The answer is: I rarely use a node object. I find it simpler to use integers, strings, or tuples as the `Location` type, and then use arrays or hash tables (dicts) that use locations as an index.

Note that the edges are *directed*: we can have an edge from A to B without also having an edge from B to A. In game maps most edges are bidirectional but sometimes there are one-way doors or jumps off cliffs that are expressed as directed edges. Let's start with an example map with both two-way and one-way links:



Part of turning a map into a graph is choosing which locations to mark. Here I decided to mark each horizontal platform as a location. We can represent this example in a graph where the `Location` type is a letter A, B, C, D, E, or F.



For each location I need a list of which locations it leads to:

```

example_graph = SimpleGraph()
example_graph.edges = {
    'A': ['B'],
    'B': ['C'],
    'C': ['B', 'D', 'F'],
    'D': ['C', 'E'],
    'E': ['F'],
    'F': [],
}

```

Before we can use it with a search algorithm, we need to make a **queue**:

```

import collections

class Queue:
    def __init__(self):
        self.elements = collections.deque()

    def empty(self) -> bool:
        return len(self.elements) == 0

    def put(self, x: T):

```

```

        self.elements.append(x)

    def get(self) -> T:
        return self.elements.popleft()

```

This queue class is a wrapper around the built-in `collections.deque` class. Feel free to use `deque` directly in your own code.

Let's try the example graph with this queue and the breadth-first search algorithm code from the main article:

```

from implementation import *

def breadth_first_search(graph: Graph, start: Location):
    # print out what we find
    frontier = Queue()
    frontier.put(start)
    reached: Dict[Location, bool] = {}
    reached[start] = True

    while not frontier.empty():
        current: Location = frontier.get()
        print(" Visiting %s" % current)
        for next in graph.neighbors(current):
            if next not in reached:
                frontier.put(next)
                reached[next] = True

    print('Reachable from A:')
    breadth_first_search(example_graph, 'A')
    print('Reachable from E:')
    breadth_first_search(example_graph, 'E')

```

```

Reachable from A:
Visiting A
Visiting B
Visiting C
Visiting D
Visiting F
Visiting E
Reachable from E:
Visiting E
Visiting F

```

Grids can be expressed as graphs too. I'll now define a new **graph** called `SquareGrid`, with `GridLocation` being a tuple `(x: int, y: int)`. In this map, the locations ("states") in the graph are the same as locations on the game map, but in many problems graph locations are not the same as map locations. Instead of storing the edges explicitly, I'll calculate them in the `neighbors` function. In many problems it's better to store them explicitly.

```
GridLocation = Tuple[int, int]

class SquareGrid:
    def __init__(self, width: int, height: int):
        self.width = width
        self.height = height
        self.walls: List[GridLocation] = []

    def in_bounds(self, id: GridLocation) -> bool:
        (x, y) = id
        return 0 <= x < self.width and 0 <= y < self.height

    def passable(self, id: GridLocation) -> bool:
        return id not in self.walls

    def neighbors(self, id: GridLocation) -> Iterator[GridLocation]:
        (x, y) = id
        neighbors = [(x+1, y), (x-1, y), (x, y-1), (x, y+1)] # E W N S
        # see "Ugly paths" section for an explanation:
        if (x + y) % 2 == 0: neighbors.reverse() # S N W E
        results = filter(self.in_bounds, neighbors)
        results = filter(self.passable, results)
        return results
```

Let's try it out with the first grid in the main article:

```
from implementation import *
g = SquareGrid(30, 15)
g.walls = DIAGRAM1_WALLS # long list, [(21, 0), (21, 2), ...]
draw_grid(g)
```

```
. . . . . . . . . . . . . . . . . . . . ##### . . . . .
. . . . . . . . . . . . . . . . . . . . ##### . . . . .
. . . . . . . . . . . . . . . . . . . . ##### . . . . .
. . . ##### . . . . . . . . . . . . . . ##### . . . . .
. . . ##### . . . . . . . . . ##### . . . . . ##### . . . . .
. . . ##### . . . . . . . . . ##### . . . . . ##### . . . . .
. . . ##### . . . . . . . . . ##### . . . . . ##### . . . . .
```

```

. . . ##### . . . . . . . . . ##### . . . . . . . . . . . . . . .
. . . ##### . . . . . . . . . ##### . . . . . . . . . . . . . . .
. . . ##### . . . . . . . . . ##### . . . . . . . . . . . . . . .
. . . ##### . . . . . . . . . ##### . . . . . . . . . . . . . . .
. . . ##### . . . . . . . . . ##### . . . . . . . . . . . . . . .
. . . . . . . . . . . . . . . ##### . . . . . . . . . . . . . . .
. . . . . . . . . . . . . . . ##### . . . . . . . . . . . . . . .
. . . . . . . . . . . . . . . ##### . . . . . . . . . . . . . . .

```

In order to reconstruct paths we need to store the location of where we came from, so I've renamed `reached` (True/False) to `came_from` (location):

```

from implementation import *

def breadth_first_search(graph: Graph, start: Location):
    frontier = Queue()
    frontier.put(start)
    came_from: Dict[Location, Optional[Location]] = {}
    came_from[start] = None

    while not frontier.empty():
        current: Location = frontier.get()
        for next in graph.neighbors(current):
            if next not in came_from:
                frontier.put(next)
                came_from[next] = current

    return came_from

g = SquareGrid(30, 15)
g.walls = DIAGRAM1_WALLS

start = (8, 7)
parents = breadth_first_search(g, start)
draw_grid(g, point_to=parents, start=start)

```

```

→ → → ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓ ← ← ← ← ← ##### ↓ ↓ ↓ ↓ ↓ ↓
→ → → → → ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓ ← ← ← ← ← ##### → ↓ ↓ ↓ ↓ ↓
→ → ↑ ##### ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓ ← ← ← ← ← ##### → → → ↓ ↓ ↓
→ ↑ ↑ ##### ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓ ← ##### ↑ ← ← ← ← ##### → → → ↓ ↓ ↓
↑ ↑ ↑ ##### → ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓ ← ##### ↑ ↑ ↑ ← ← ##### ↓ ↓ ↓ ↓
↑ ↑ ↑ ##### → → ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓ ← ##### ↑ ↑ ↑ ↑ ↑ ← ← ##### ↓ ↓ ↓ ↓
↓ ↓ ↓ ##### → → → A ← ← ← ← ← ##### ↑ ↑ ↑ ↑ ↑ ← ← ← ← ← ← ← ← ←
↓ ↓ ↓ ##### → → ↑ ↑ ↑ ↑ ↑ ↑ ↑ ↑ ↑ ← ← ##### ↑ ↑ ↑ ↑ ↑ ← ← ← ← ←
↓ ↓ ↓ ##### → ↑ ↑ ↑ ↑ ↑ ↑ ↑ ↑ ↑ ← ← ##### ↑ ↑ ↑ ↑ ↑ ← ← ← ← ←
→ ↓ ↓ ##### ↑ ↑ ↑ ↑ ↑ ↑ ↑ ↑ ↑ ##### ↑ ↑ ↑ ↑ ↑ ← ← ← ← ← ← ← ← ←
→ → ↓ ##### ↑ ↑ ↑ ↑ ↑ ↑ ↑ ↑ ↑ ##### ↑ ↑ ↑ ↑ ↑ ↑ ↑ ↑ ↑ ← ← ← ← ← ←
→ → → → → ↑ ↑ ↑ ↑ ↑ ↑ ↑ ↑ ↑ ##### ↑ ↑ ↑ ↑ ↑ ↑ ↑ ↑ ↑ ↑ ↑ ← ← ← ← ←

```

Some implementations use *internal storage*, creating a Node object to hold `came_from` and other values for each graph node. I've instead chosen to use *external storage*, creating a single hash table to store the `came_from` for all graph nodes. If you know your map locations have integer indices, another option is to use an array to store `came_from`.

1.2 Early Exit

#

Following the code from the main article, we need to add an *if* statement to the main loop. This test is optional for Breadth First Search or Dijkstra's Algorithm and effectively required for Greedy Best-First Search and A*:

```
from implementation import *

def breadth_first_search(graph: Graph, start: Location, goal: Location):
    frontier = Queue()
    frontier.put(start)
    came_from: Dict[Location, Optional[Location]] = {}
    came_from[start] = None

    while not frontier.empty():
        current: Location = frontier.get()


        if current == goal:
            break

        for next in graph.neighbors(current):
            if next not in came_from:
                frontier.put(next)
                came_from[next] = current

    return came_from

g = SquareGrid(30, 15)
g.walls = DIAGRAM1_WALLS

start = (8, 7)
goal = (17, 2)
parents = breadth_first_search(g, start, goal)
draw_grid(g, point_to=parents, start=start, goal=goal)
```



You can see that the algorithm stops when it finds the goal z .

1.3 Dijkstra's Algorithm

#

This is what adds complexity to graph search, because we're going to start processing locations in a better order than "first in, first out". What do we need to change?

1. The *graph* needs to know cost of movement.
2. The *queue* needs to return nodes in a different order.
3. The *search* needs to keep track of these costs from the graph and give them to the queue.

1.3.1 Graph with weights

A regular graph tells me the `neighbors` of each node. A *weighted* graph also tells me the cost of moving along each edge. I'm going to add a `cost(from_node, to_node)` function that tells us the cost of moving from location `from_node` to its neighbor `to_node`. Here's the interface:

```
class WeightedGraph(Graph):
    def cost(self, from_id: Location, to_id: Location) -> float: pass
```

Let's implement the interface with a grid that uses grid locations and stores the weights in a dict:

```

class GridWithWeights(SquareGrid):
    def __init__(self, width: int, height: int):
        super().__init__(width, height)
        self.weights: Dict[GridLocation, float] = {}

    def cost(self, from_node: GridLocation, to_node: GridLocation) -> float:
        return self.weights.get(to_node, 1)

```

In this forest map I chose to make movement depend only on `to_node`, but there are other types of movement that use both nodes^[1]. An alternate implementation would be to include the movement costs in the value returned by the `neighbors` function.

1.3.2 Queue with priorities

A priority queue associates with each item a number called a “priority”. When returning an item, it picks the one with the lowest number.

insert

Add item to queue

remove

Remove item with the lowest number

reprioritize

(optional) Change an existing item’s priority to a lower number

Here’s a reasonably fast priority queue that uses *binary heaps*, but does not support reprioritize. To get the right ordering, we’ll use tuples (priority, item). When an element is inserted that is already in the queue, we’ll have a duplicate; I’ll explain why that’s ok in the Optimization section.

```

import heapq

class PriorityQueue:
    def __init__(self):
        self.elements: List[Tuple[float, T]] = []

    def empty(self) -> bool:

```

```

    return len(self.elements) == 0

    def put(self, item: T, priority: float):
        heapq.heappush(self.elements, (priority, item))

    def get(self) -> T:
        return heapq.heappop(self.elements)[1]

```

1.3.3 Search

Here's a tricky bit about the implementation: once we add movement costs it's possible to visit a location again, with a better `cost_so_far`. That means the line `if next not in came_from` won't work. Instead, have to check if the cost has gone down since the last time we reached. (In the original version of the article I wasn't checking this, but my code worked anyway; [I wrote some notes about that bug.](#))

This forest map is from [the main page](#).

```

def dijkstra_search(graph: WeightedGraph, start: Location, goal: Location):
    frontier = PriorityQueue()
    frontier.put(start, 0)
    came_from: Dict[Location, Optional[Location]] = {}
    cost_so_far: Dict[Location, float] = {}
    came_from[start] = None
    cost_so_far[start] = 0

    while not frontier.empty():
        current: Location = frontier.get()

        if current == goal:
            break

        for next in graph.neighbors(current):
            new_cost = cost_so_far[current] + graph.cost(current, next)
            if next not in cost_so_far or new_cost < cost_so_far[next]:
                cost_so_far[next] = new_cost
                priority = new_cost
                frontier.put(next, priority)
                came_from[next] = current

    return came_from, cost_so_far

```

Finally, after searching I need to build the path:

```
def reconstruct_path(came_from: Dict[Location, Location],
                    start: Location, goal: Location) -> List[Location]:
    current: Location = goal
    path: List[Location] = []
    while current != start:
        path.append(current)
        current = came_from[current]
    path.append(start) # optional
    path.reverse() # optional
    return path
```

Although paths are best thought of as a sequence of edges, it's convenient to store them as a sequence of nodes. To build the path, start at the end and follow the `came_from` map, which points to the previous node. When we reach start, we're done. It is the **backwards** path, so call `reverse()` at the end of `reconstruct_path` if you need it to be stored forwards. Sometimes it's actually more convenient to store it backwards. Sometimes it's useful to also store the start node in the list.

Let's try it out:

```
from implementation import *
start, goal = (1, 4), (8, 3)
came_from, cost_so_far = dijkstra_search(diagram4, start, goal)
draw_grid(diagram4, point_to=came_from, start=start, goal=goal)
print()
draw_grid(diagram4, path=reconstruct_path(came_from, start=start, goal=goal))
```

```
↓ ↓ ← ← ← ← ← ← ←
↓ ↓ ← ← ← ↑ ↑ ← ← ←
↓ ↓ ← ← ← ← ↑ ↑ ← ←
↓ ↓ ← ← ← ← ↑ Z .
→ A ← ← ← ← . . .
↑ ↑ ← ← ← ← . . .
↑ ↑ ← ← ← ← . . .
↑ ##### ↑ ← ↓ ↓ . .
↑ ##### ↓ ↓ ↓ ← ← .
↑ ← ← ← ← ← ← ← .
. @ @ @ @ @ @ . . .
```

```

. @ . . . . @ @ . .
. @ . . . . . @ @ .
. @ . . . . . . @ .
. @ . . . . . . . .
. . . . . . . . . .
. . . . . . . . . .
. ##### . . . . .
. ##### . . . . .
. . . . . . . . . .

```

The first output shows the vector field; the second shows the path.

Why is the path going up and over? Remember that this is the forest example from the main page, where the middle of the map has a big forest that's slow to move through. The shortest path goes around the forest, not through it.

The line `if next not in cost_so_far or new_cost < cost_so_far[next]` could be simplified to `if new_cost < cost_so_far.get(next, Infinity)` but I didn't want to explain Python's `get()` in the main article so I left it as is. Another approach would be to use `collections.defaultdict` defaulting to infinity.

Collecting distances instead of directions gives us a distance field, which can be useful for some applications:

```

from implementation import *
start, goal = (1, 4), (8, 3)
came_from, cost_so_far = dijkstra_search(diagram4, start, goal)
draw_grid(diagram4, number=cost_so_far, start=start, goal=goal)

```

```

5 4 5 6 7 8 9 10 11 12
4 3 4 5 10 13 10 11 12 13
3 2 3 4 9 14 15 12 13 14
2 1 2 3 8 13 18 17 2 .
1 A 1 6 11 16 . . . .
2 1 2 7 12 17 . . . .
3 2 3 4 9 14 19 . . .
4 ##### 14 19 18 15 . .
5 ##### 15 16 13 14 15 .
6 7 8 9 10 11 12 13 14 .

```

1.4 A* Search

#

A* is almost exactly like Dijkstra's Algorithm, except we add in a heuristic. Note that the code for the algorithm *isn't specific to grids*. Knowledge about grids is in the graph class (`GridWithWeights`), the locations, and in the `heuristic` function. Replace those three and you can use the A* algorithm code with any other graph structure.

```
def heuristic(a: GridLocation, b: GridLocation) -> float:
    (x1, y1) = a
    (x2, y2) = b
    return abs(x1 - x2) + abs(y1 - y2)

def a_star_search(graph: WeightedGraph, start: Location, goal: Location):
    frontier = PriorityQueue()
    frontier.put(start, 0)
    came_from: Dict[Location, Optional[Location]] = {}
    cost_so_far: Dict[Location, float] = {}
    came_from[start] = None
    cost_so_far[start] = 0

    while not frontier.empty():
        current: Location = frontier.get()

        if current == goal:
            break

        for next in graph.neighbors(current):
            new_cost = cost_so_far[current] + graph.cost(current, next)
            if next not in cost_so_far or new_cost < cost_so_far[next]:
                cost_so_far[next] = new_cost
                priority = new_cost + heuristic(next, goal)
                frontier.put(next, priority)
                came_from[next] = current

    return came_from, cost_so_far
```

Let's try it out:

```
from implementation import *
start, goal = (1, 4), (8, 3)
came_from, cost_so_far = a_star_search(diagram4, start, goal)
draw_grid(diagram4, point_to=came_from, start=start, goal=goal)
print()
draw_grid(diagram4, path=reconstruct_path(came_from, start=start, goal=goal))
```

```

↓ ↓ ↓ ↓ ← ← ← ← ← ←
↓ ↓ ↓ ↓ ← ↑ ↑ ← ← ←
↓ ↓ ↓ ↓ ← ← ↑ ↑ ← ←
→ ↓ ← ← ← ← . ↑ Z .
→ A ← ← ← . . . . .
↑ ↑ ↑ ← ← . . . . .
↑ ↑ ↑ ← ← . . . . .
↑ ##### . . . . .
. ##### . . . . .
. . . . . . . . . .

. . . @ @ @ @ . . .
. . . @ . . @ @ . .
. . . @ . . . @ @ .
. @ @ @ . . . @ .
. @ . . . . . . .
. . . . . . . . .
. . . . . . . . .
. ##### . . . . .
. ##### . . . . .
. . . . . . . . .

```

Here are the distances it calculated:

```

from implementation import *
start, goal = (1, 4), (8, 3)
came_from, cost_so_far = a_star_search(diagram4, start, goal)
draw_grid(diagram4, number=cost_so_far, start=start, goal=goal)

```

```

5  4  5  6  7  8  9 10 11 12
4  3  4  5 10 13 10 11 12 13
3  2  3  4  9 14 15 12 13 14
2  1  2  3  8 13 . 17 Z .
1  A  1  6 11 . . . . .
2  1  2  7 12 . . . . .
3  2  3  4  9 . . . . .
4 ##### . . . . .
. ##### . . . . .
. . . . . . . . .

```

And that's it! We've implemented graphs, grids, Breadth First Search, Dijkstra's Algorithm, and A*.

1.4.1 Straighter paths

If you implement this code in your own project you might find that some of the paths aren't as "straight" as you'd like. **This is normal**. When using *grids*, especially grids where every step has the same movement cost, you end up with **ties**: many paths have exactly the same cost. A* ends up picking one of the many short paths, and often **it won't look good to you**. I list [some solutions](#) in a later section.

2 C++ Implementation

#

Note: some of the sample code needs to include [redblobgames/pathfinding/a-star/implementation.cpp](#) to run. I am using C++14 for this code so some of it will need to be changed if you use an older version of the C++ standard.

The code here is meant for the tutorial and is not production-quality; there's a section at the end with tips on making it better.

2.1 Breadth First Search

#

Let's implement Breadth First Search in C++. These are the components we need:

Graph

a data structure that can tell me the `neighbors` for each graph location (see [this tutorial](#)). A *weighted* graph can also tell me the `cost` of moving along an edge.

Locations

a simple value (int, string, tuple, etc.) that *labels* locations in the graph. These are not necessarily locations on the map. They may include additional information such as direction, fuel, lane, or inventory, depending on the problem being solved.

Search

an algorithm that takes a graph, a starting graph location, and optionally a goal graph location, and calculates some useful information (reached,

parent pointer, distance) for some or all graph locations.


Queue

a data structure used by the search algorithm to decide the order in which to process the graph locations.


In the main article, I focused on **search**. On this page, I'll fill in the rest of the details to make complete working programs. Let's start with a **graph** where the locations are `char`:

```
struct SimpleGraph {  
    std::unordered_map<char, std::vector<char> > edges;  
  
    std::vector<char> neighbors(char id) {  
        return edges[id];  
    }  
};
```

Note that the edges are *directed*: we can have an edge from A to B without also having an edge from B to A. In game maps most edges are bidirectional but sometimes there are one-way doors or jumps off cliffs that are expressed as directed edges. Let's start with an example map with both two-way and one-way links:



Part of turning a map into a graph is choosing which locations to mark. Here I decided to mark each horizontal platform as a location. We can represent this example in a graph where the `Location` type is a letter A, B, C, D, E, or F.



```

SimpleGraph example_graph {{
    {'A', {'B'}},
    {'B', {'C'}},
    {'C', {'B', 'D', 'F'}},
    {'D', {'C', 'E'}},
    {'E', {'F'}},
    {'F', {}},
  }};

```

The C++ standard library already includes a queue class. We now have a graph (`SimpleGraph`), locations (`char`), and a queue (`std::queue`). Now we can try Breadth First Search:

```

#include "redblobgames/pathfinding/a-star/implementation.cpp"

void breadth_first_search(SimpleGraph graph, char start) {
    std::queue<char> frontier;
    frontier.push(start);

    std::unordered_set<char> reached;
    reached.insert(start);

    while (!frontier.empty()) {
        char current = frontier.front();

```

```

    frontier.pop();

    std::cout << "  Visiting " << current << '\n';
    for (char next : graph.neighbors(current)) {
        if (reached.find(next) == reached.end()) {
            frontier.push(next);
            reached.insert(next);
        }
    }
}
}

int main() {
    std::cout << "Reachable from A:\n";
    breadth_first_search(example_graph, 'A');
    std::cout << "Reachable from E:\n";
    breadth_first_search(example_graph, 'E');
}

```

```

Reachable from A:
  Visiting A
  Visiting B
  Visiting C
  Visiting D
  Visiting F
  Visiting E
Reachable from E:
  Visiting E
  Visiting F

```

Grids can be expressed as graphs too. I'll now define a new **graph** called `SquareGrid`, with **locations** structs with two ints. In this map, the locations ("states") in the graph are the same as locations on the game map, but in many problems graph locations are not the same as map locations. Instead of storing the edges explicitly, I'll calculate them in the `neighbors` function. In many problems it's better to store them explicitly.

```

struct GridLocation {
    int x, y;
};

namespace std {
    /* implement hash function so we can put GridLocation into an unordered_set */
    template <> struct hash<GridLocation> {
        typedef GridLocation argument_type;
        typedef std::size_t result_type;
    };
}

```

```

    std::size_t operator()(const GridLocation& id) const noexcept {
        return std::hash<int>()(id.x ^ (id.y << 4));
    }
};

}

struct SquareGrid {
    static std::array<GridLocation, 4> DIRS;

    int width, height;
    std::unordered_set<GridLocation> walls;

    SquareGrid(int width_, int height_)
        : width(width_), height(height_) {}

    bool in_bounds(GridLocation id) const {
        return 0 <= id.x && id.x < width
            && 0 <= id.y && id.y < height;
    }

    bool passable(GridLocation id) const {
        return walls.find(id) == walls.end();
    }

    std::vector<GridLocation> neighbors(GridLocation id) const {
        std::vector<GridLocation> results;

        for (GridLocation dir : DIRS) {
            GridLocation next{id.x + dir.x, id.y + dir.y};
            if (in_bounds(next) && passable(next)) {
                results.push_back(next);
            }
        }

        if ((id.x + id.y) % 2 == 0) {
            // see "Ugly paths" section for an explanation:
            std::reverse(results.begin(), results.end());
        }

        return results;
    }
};

std::array<GridLocation, 4> SquareGrid::DIRS = {
    /* East, West, North, South */
    GridLocation{1, 0}, GridLocation{-1, 0},

```



```

        frontier.push(next);
        came_from[next] = current;
    }
}
}
return came_from;
}

int main() {
    SquareGrid grid = make_diagram1();
    GridLocation start{7, 8};
    auto parents = breadth_first_search(grid, start);
    draw_grid(grid, nullptr, &parents, nullptr, &start);
}

```

Some implementations use *internal storage*, creating a Node object to hold `came_from` and other values for each graph node. I've instead chosen to use *external storage*, creating a single `std::unordered_map` to store the `came_from` for all graph nodes. If you know your map locations have integer indices, another option is to use a 1D or 2D array/vector to store `came_from` and other values.

2.2 Early Exit

#

Breadth First Search and Dijkstra's Algorithm will explore the entire map by default. If we're looking for a path to a single point we can add `if (current == goal)` to exit the loop as soon as we find the path.

```

#include "redblobgames/pathfinding/a-star/implementation.cpp"

template<typename Location, typename Graph>
std::unordered_map<Location, Location>
breadth_first_search(Graph graph, Location start, Location goal) {
    std::queue<Location> frontier;
    frontier.push(start);

    std::unordered_map<Location, Location> came_from;
    came_from[start] = start;

    while (!frontier.empty()) {
        Location current = frontier.front();
        frontier.pop();

        if (current == goal) {
            break;
        }

        for (Location next : graph.neighbors(current)) {
            if (came_from.find(next) == came_from.end()) {
                frontier.push(next);
                came_from[next] = current;
            }
        }
    }
    return came_from;
}

int main() {
    GridLocation start{8, 7}, goal{17, 2};
    SquareGrid grid = make_diagram1();
    auto came_from = breadth_first_search(grid, start, goal);
    draw_grid(grid, nullptr, &came_from, nullptr, &start, &goal);
}

```

```

. → → ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓ ← . . . ##### . . . . .
→ → → → ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓ ← ← ← . . . ##### . . . . .
→ → → → → ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓ ← ← ← Z . . . ##### . . . . .
→ → ↑ ##### ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓ ← ← ← ← ← ← . . . ##### . . . . .
. ↑ ↑ ##### ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓ ← ##### ↑ ← ← . . . ##### . . . . .
. . ↑ ##### → ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓ ← ← ##### ↑ ↑ . . . ##### . . . . .
. . . ##### → → ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓ ← ← ##### ↑ . . . ##### . . . . .
. . . ##### → → → A ← ← ← ← ← ← ##### . . . . . ##### . . . . .
. . . ##### → → ↑ ↑ ↑ ← ← ← ##### . . . . . ##### . . . . .
. . ↓ ##### → ↑ ↑ ↑ ↑ ↑ ↑ ← ← ##### . . . . . ##### . . . . .
. ↓ ↓ ##### ↑ ↑ ↑ ↑ ↑ ↑ ↑ ↑ ↑ ← ##### . . . . . ##### . . . . .
→ → ↓ ##### ↑ ↑ ↑ ↑ ↑ ↑ ↑ ↑ ↑ ##### . . . . . ##### . . . . .
→ → → → → ↑ ↑ ↑ ↑ ↑ ↑ ↑ ↑ ↑ ##### . . . . . ##### . . . . .
→ → → → → ↑ ↑ ↑ ↑ ↑ ↑ ↑ ↑ ↑ ##### . . . . . ##### . . . . .
. → → ↑ ↑ ↑ ↑ ↑ ↑ ↑ ↑ ↑ ##### . . . . . ##### . . . . .

```

In the output we can see that the algorithm did not explore the entire map, but stopped early.

2.3 Dijkstra's Algorithm

#

This is what adds complexity to graph search, because we're going to start processing locations in a better order than "first in, first out". What do we need to change?

1. The *graph* needs to know cost of movement.
2. The *queue* needs to return nodes in a different order.
3. The *search* needs to keep track of these costs from the graph and give them to the queue.

2.3.1 Graph with weights

A regular graph tells me the `neighbors` of each node. A *weighted* graph also tells me the cost of moving along each edge. I'm going to add a `cost(from_node, to_node)` function that tells us the cost of moving from location `from_node` to its neighbor `to_node`. In this forest map I chose to make movement depend only on `to_node`, but there are other types of movement that use both nodes^[2]. An alternate implementation would be to merge this into the `neighbors` function. Here's a grid with a list of forest tiles, which will have movement cost 5:

```
struct GridWithWeights: SquareGrid {
    std::unordered_set<GridLocation> forests;
    GridWithWeights(int w, int h): SquareGrid(w, h) {}
    double cost(GridLocation from_node, GridLocation to_node) const {
        return forests.find(to_node) != forests.end()? 5 : 1;
    }
};
```

2.3.2 Queue with priorities

We need a priority queue. C++ offers a `priority_queue` class that uses a binary heap but not the reprioritize operation. I'll use a pair (priority, item) for the queue elements to get the right ordering. By default, the C++ priority queue returns the maximum element first, using the `std::less` comparator; we want the minimum element instead, so I'll use the `std::greater` comparator.

```
template<typename T, typename priority_t>
struct PriorityQueue {
    typedef std::pair<priority_t, T> PQElement;
    std::priority_queue<PQElement, std::vector<PQElement>,
                      std::greater<PQElement>> elements;

    inline bool empty() const {
        return elements.empty();
    }

    inline void put(T item, priority_t priority) {
        elements.emplace(priority, item);
    }

    T get() {
        T best_item = elements.top().second;
        elements.pop();
        return best_item;
    }
};
```

In this sample code I'm wrapping the C++ `std::priority_queue` class but I think it'd be reasonable to use that class directly without the wrapper.

2.3.3 Search

See [the forest map from the main page](#).

```
template<typename Location, typename Graph>
void dijkstra_search
    (Graph graph,
     Location start,
     Location goal,
```

```

    std::unordered_map<Location, Location>& came_from,
    std::unordered_map<Location, double>& cost_so_far)
{
    PriorityQueue<Location, double> frontier;
    frontier.put(start, 0);

    came_from[start] = start;
    cost_so_far[start] = 0;

    while (!frontier.empty()) {
        Location current = frontier.get();

        if (current == goal) {
            break;
        }

        for (Location next : graph.neighbors(current)) {
            double new_cost = cost_so_far[current] + graph.cost(current, next);
            if (cost_so_far.find(next) == cost_so_far.end()
                || new_cost < cost_so_far[next]) {
                cost_so_far[next] = new_cost;
                came_from[next] = current;
                frontier.put(next, new_cost);
            }
        }
    }
}

```

The types of the `cost` variables should all match the types used in the graph. If you use `int` then you can use `int` for the cost variable and the priorities in the priority queue; if you use `double` then you should use `double` for these. In this code I used `double` but I could've used `int` and it would've worked the same. However, if your graph edge costs are doubles or if your heuristic uses doubles, then you'll need to use doubles here.

Finally, after searching I need to build the path:

```

template<typename Location>
std::vector<Location> reconstruct_path(
    Location start, Location goal,
    std::unordered_map<Location, Location> came_from
) {
    std::vector<Location> path;

```

```

Location current = goal;
while (current != start) {
    path.push_back(current);
    current = came_from[current];
}
path.push_back(start); // optional
std::reverse(path.begin(), path.end());
return path;
}

```

Although paths are best thought of as a sequence of edges, it's convenient to store them as a sequence of nodes. To build the path, start at the end and follow the `came_from` map, which points to the previous node. When we reach start, we're done. It is the **backwards** path, so call `reverse()` at the end of `reconstruct_path` if you need it to be stored forwards. Sometimes it's actually more convenient to store it backwards. Sometimes it's useful to also store the start node in the list.

Let's try it out:

```

#include "redblobgames/pathfinding/a-star/implementation.cpp"

int main() {
    GridWithWeights grid = make_diagram4();
    GridLocation start{1, 4}, goal{8, 3};
    std::unordered_map<GridLocation, GridLocation> came_from;
    std::unordered_map<GridLocation, double> cost_so_far;
    dijkstra_search(grid, start, goal, came_from, cost_so_far);
    draw_grid(grid, nullptr, &came_from, nullptr, &start, &goal);
    std::cout << '\n';
    std::vector<GridLocation> path = reconstruct_path(start, goal, came_from);
    draw_grid(grid, nullptr, nullptr, &path, &start, &goal);
    std::cout << '\n';
    draw_grid(grid, &cost_so_far, nullptr, nullptr, &start, &goal);
}

```

```

↓ ↓ ← ← ← ← ← ← ← ←
↓ ↓ ← ← ← ↑ ↑ ← ← ←
↓ ↓ ← ← ← ← ↑ ↑ ← ←
↓ ↓ ← ← ← ← ← ↑ Z .
→ A ← ← ← ← . . . .
↑ ↑ ← ← ← ← . . . .
↑ ↑ ← ← ← ← ← . . .
↑ ##### ↑ ← ↓ ↓ . .
↑ ##### ↓ ↓ ↓ ← ← .

```

```

↑ ← ← ← ← ← ← ← .
. @ @ @ @ @ @ . . .
. @ . . . . @ @ . .
. @ . . . . . @ @ .
. @ . . . . . . Z .
. A . . . . . . . .
. . . . . . . . .
. . . . . . . . .
. ##### . . . . .
. ##### . . . . .
. . . . . . . . .

5 4 5 6 7 8 9 10 11 12
4 3 4 5 10 13 10 11 12 13
3 2 3 4 9 14 15 12 13 14
2 1 2 3 8 13 18 17 Z .
1 A 1 6 11 16 . . . .
2 1 2 7 12 17 . . . .
3 2 3 4 9 14 19 . . .
4 ##### 14 19 18 15 . .
5 ##### 15 16 13 14 15 .
6 7 8 9 10 11 12 13 14 .

```

Why is the path going up and over? Remember that this is the forest example from the main page, where the middle of the map has a big forest that's slow to move through. The shortest path goes around the forest, not through it.

The results are not always the same as the Python version because I'm using the built-in priority queues in C++ and Python. These may order equal-valued nodes differently. **This is something you'll run into if using grids.** There are *many* equally short paths, and the pathfinder will find *one* of them, not necessarily the one that looks the best to your eye.

2.4 A* Search

#

A* is almost exactly like Dijkstra's Algorithm, except we add in a heuristic. Note that the code for the algorithm *isn't specific to grids*. Knowledge about grids is in the graph class (`GridWithWeights`), the locations (`Location` struct), and in the `heuristic` function. Replace those three and you can use the A* algorithm code with any other graph structure.

```

inline double heuristic(GridLocation a, GridLocation b) {
    return std::abs(a.x - b.x) + std::abs(a.y - b.y);
}

```

```

template<typename Location, typename Graph>
void a_star_search
    (Graph graph,
     Location start,
     Location goal,
     std::unordered_map<Location, Location>& came_from,
     std::unordered_map<Location, double>& cost_so_far)
{
    PriorityQueue<Location, double> frontier;
    frontier.put(start, 0);

    came_from[start] = start;
    cost_so_far[start] = 0;

    while (!frontier.empty()) {
        Location current = frontier.get();

        if (current == goal) {
            break;
        }

        for (Location next : graph.neighbors(current)) {
            double new_cost = cost_so_far[current] + graph.cost(current, next);
            if (cost_so_far.find(next) == cost_so_far.end()
                || new_cost < cost_so_far[next]) {
                cost_so_far[next] = new_cost;
                double priority = new_cost + heuristic(next, goal);
                frontier.put(next, priority);
                came_from[next] = current;
            }
        }
    }
}

```

The type of the `priority` values including the type used in the priority queue should be big enough to include both the graph costs (`cost_t`) and the heuristic value. For example, if the graph costs are ints and the heuristic returns a double, then you need the priority queue to accept doubles. In this sample code I use `double` for all three (cost, heuristic, and priority), but I could've used `int` because my costs and heuristics are integer valued.

Minor note: It would be more correct to write `frontier.put(start, heuristic(start, goal))` than `frontier.put(start, 0)` but it makes no difference here because the start node's priority doesn't matter. It is the only node in the priority queue and it is selected and removed before anything else is put in there.

Let's try it out:

```
#include "redblobgames/pathfinding/a-star/implementation.cpp"

int main() {
    GridWithWeights grid = make_diagram4();
    GridLocation start{1, 4}, goal{8, 3};
    std::unordered_map<GridLocation, GridLocation> came_from;
    std::unordered_map<GridLocation, double> cost_so_far;
    a_star_search(grid, start, goal, came_from, cost_so_far);
    draw_grid(grid, nullptr, &came_from, nullptr, &start, &goal);
    std::cout << '\n';
    std::vector<GridLocation> path = reconstruct_path(start, goal, came_from);
    draw_grid(grid, nullptr, nullptr, &path, &start, &goal);
    std::cout << '\n';
    draw_grid(grid, &cost_so_far, nullptr, nullptr, &start, &goal);
}
```

```
↓ ↓ ↓ ↓ ← ← ← ← ←
↓ ↓ ↓ ↓ ← ↑ ↑ ← ← ←
↓ ↓ ↓ ↓ ← ← ↑ ↑ ← ←
→ ↓ ← ← ← ← . ↑ Z .
→ A ← ← ← . . . . .
↑ ↑ ↑ ← ← . . . . .
↑ ↑ ↑ ← ← . . . . .
↑ ##### . . . . .
. ##### . . . . .
. . . . . . . . . .

. . . @ @ @ @ . . .
. . . @ . . @ @ . .
. . . @ . . . @ @ .
. @ @ @ . . . . Z .
. A . . . . . . .
. . . . . . . . .
. . . . . . . . .
. ##### . . . . .
. ##### . . . . .
. . . . . . . . .

5 4 5 6 7 8 9 10 11 12
4 3 4 5 10 13 10 11 12 13
3 2 3 4 9 14 15 12 13 14
2 1 2 3 8 13 . 17 Z .
1 A 1 6 11 . . . . .
2 1 2 7 12 . . . . .
```

```

3  2  3  4  9  .  .  .  .  .
4  #####  .  .  .  .  .
.  #####  .  .  .  .  .
.  .  .  .  .  .  .  .  .

```

And that's it! We've implemented graphs, grids, Breadth First Search, Dijkstra's Algorithm, and A*.

2.4.1 Straighter paths

If you implement this code in your own project you might find that some of the paths aren't as "straight" as you'd like. **This is normal**. When using *grids*, especially grids where every step has the same movement cost, you end up with **ties**: many paths have exactly the same cost. A* ends up picking one of the many short paths, and often **it won't look good to you**. I list [some solutions](#) in a later section.

2.5 Production code

#

The C++ code I've shown above is simplified to make it easier to follow the algorithm and data structures. In practice there are many things you'd want to do differently:

- inlining small functions
- the `Location` parameter should be part of the `Graph`
- the cost could be `int` or `double`, and should be part of the `Graph`
- use `array` instead of `unordered_set` if the ids are dense integers, and reset these values on exit instead of initializing on entry
- pass larger data structures by reference instead of by value
- return larger data structures in out parameters instead of returning them, or use move constructors (for example, the vector returned from the `neighbors` function)
- the heuristic can vary and should be a template parameter to the A* function so that it can be inlined

Here's how the A* code might look different with some (but not all) of these changes:

```

template<typename Graph>
void a_star_search
(Graph graph,
 typename Graph::Location start,
 typename Graph::Location goal,
 std::function<typename Graph::cost_t(typename Graph::Location a,
                                     typename Graph::Location b)> heuristic,
 std::unordered_map<typename Graph::Location,
                   typename Graph::Location>& came_from,
 std::unordered_map<typename Graph::Location,
                   typename Graph::cost_t>& cost_so_far)
{
    typedef typename Graph::Location Location;
    typedef typename Graph::cost_t cost_t;
    PriorityQueue<Location, cost_t> frontier;
    std::vector<Location> neighbors;
    frontier.put(start, cost_t(0));

    came_from[start] = start;
    cost_so_far[start] = cost_t(0);

    while (!frontier.empty()) {
        typename Location current = frontier.get();

        if (current == goal) {
            break;
        }

        graph.get_neighbors(current, neighbors);
        for (Location next : neighbors) {
            cost_t new_cost = cost_so_far[current] + graph.cost(current, next);
            if (cost_so_far.find(next) == cost_so_far.end()
                || new_cost < cost_so_far[next]) {
                cost_so_far[next] = new_cost;
                cost_t priority = new_cost + heuristic(next, goal);
                frontier.put(next, priority);
                came_from[next] = current;
            }
        }
    }
}

```

I wanted the code on this page to be about the algorithms and data structures and not about the C++ optimizations so I tried to show simple code instead of fast or abstract code.

3 C# Implementation

#

These were my first C# programs so they might not be idiomatic or stylistically proper. These examples aren't as complete as the Python and C++ sections, but I hope they're helpful.

Here's a simple graph, and Breadth First Search:

```
using System;
using System.Collections.Generic;

public class Graph<Location>
{
    // NameValueCollection would be a reasonable alternative here, if
    // you're always using string location types
    public Dictionary<Location, Location[]> edges
        = new Dictionary<Location, Location[]>();

    public Location[] Neighbors(Location id)
    {
        return edges[id];
    }
};

class BreadthFirstSearch
{
    static void Search(Graph<string> graph, string start)
    {
        var frontier = new Queue<string>();
        frontier.Enqueue(start);

        var reached = new HashSet<string>();
        reached.Add(start);

        while (frontier.Count > 0)
        {
            var current = frontier.Dequeue();

            Console.WriteLine("Visiting {0}", current);
            foreach (var next in graph.Neighbors(current))
            {
```

```

        if (!reached.Contains(next)) {
            frontier.Enqueue(next);
            reached.Add(next);
        }
    }
}

static void Main()
{
    Graph<string> g = new Graph<string>();
    g.edges = new Dictionary<string, string[]>
    {
        { "A", new [] { "B" } },
        { "B", new [] { "A", "C", "D" } },
        { "C", new [] { "A" } },
        { "D", new [] { "E", "A" } },
        { "E", new [] { "B" } }
    };

    Search(g, "A");
}

```

Here's a graph representing a grid with weighted edges (the forest and walls example from the main page):

```

using System;
using System.Collections.Generic;

// A* needs only a WeightedGraph and a location type L, and does *not*
// have to be a grid. However, in the example code I am using a grid.
public interface WeightedGraph<L>
{
    double Cost(Location a, Location b);
    IEnumerable<Location> Neighbors(Location id);
}

public struct Location
{
    // Implementation notes: I am using the default Equals but it can
    // be slow. You'll probably want to override both Equals and
    // GetHashCode in a real project.

```

```
public readonly int x, y;
public Location(int x, int y)
{
    this.x = x;
    this.y = y;
}
}

public class SquareGrid : WeightedGraph<Location>
{
    // Implementation notes: I made the fields public for convenience,
    // but in a real project you'll probably want to follow standard
    // style and make them private.

    public static readonly Location[] DIRS = new []
    {
        new Location(1, 0),
        new Location(0, -1),
        new Location(-1, 0),
        new Location(0, 1)
    };

    public int width, height;
    public HashSet<Location> walls = new HashSet<Location>();
    public HashSet<Location> forests = new HashSet<Location>();

    public SquareGrid(int width, int height)
    {
        this.width = width;
        this.height = height;
    }

    public bool InBounds(Location id)
    {
        return 0 <= id.x && id.x < width
            && 0 <= id.y && id.y < height;
    }

    public bool Passable(Location id)
    {
        return !walls.Contains(id);
    }

    public double Cost(Location a, Location b)
    {
        return forests.Contains(b) ? 5 : 1;
    }
}
```

```

public IEnumerable<Location> Neighbors(Location id)
{
    foreach (var dir in DIRS) {
        Location next = new Location(id.x + dir.x, id.y + dir.y);
        if (InBounds(next) && Passable(next)) {
            yield return next;
        }
    }
}

public class PriorityQueue<T>
{
    // I'm using an unsorted array for this example, but ideally this
    // would be a binary heap. There's an open issue for adding a binary
    // heap to the standard C# library: https://github.com/dotnet/corefx/issues/11131
    // Until then, find a binary heap class:
    // * https://github.com/BlueRaja/High-Speed-Priority-Queue-for-C-Sharp
    // * http://visualstudiomagazine.com/articles/2012/11/01/priority-queue-in-net.aspx
    // * http://xfleury.github.io/graphsearch.html
    // * http://stackoverflow.com/questions/102398/priority-queue-in-net

    private List<Tuple<T, double>> elements = new List<Tuple<T, double>>();

    public int Count
    {
        get { return elements.Count; }
    }

    public void Enqueue(T item, double priority)
    {
        elements.Add(Tuple.Create(item, priority));
    }

    public T Dequeue()
    {
        int bestIndex = 0;

        for (int i = 0; i < elements.Count; i++) {
            if (elements[i].Item2 < elements[bestIndex].Item2) {
                bestIndex = i;
            }
        }

        T bestItem = elements[bestIndex].Item1;
    }
}

```

```

        elements.RemoveAt(bestIndex);
        return bestItem;
    }
}

/* NOTE about types: in the main article, in the Python code I just
 * use numbers for costs, heuristics, and priorities. In the C++ code
 * I use a typedef for this, because you might want int or double or
 * another type. In this C# code I use double for costs, heuristics,
 * and priorities. You can use an int if you know your values are
 * always integers, and you can use a smaller size number if you know
 * the values are always small. */

public class AStarSearch
{
    public Dictionary<Location, Location> cameFrom
        = new Dictionary<Location, Location>();
    public Dictionary<Location, double> costSoFar
        = new Dictionary<Location, double>();

    // Note: a generic version of A* would abstract over Location and
    // also Heuristic
    static public double Heuristic(Location a, Location b)
    {
        return Math.Abs(a.x - b.x) + Math.Abs(a.y - b.y);
    }

    public AStarSearch(WeightedGraph<Location> graph, Location start, Location goal)
    {
        var frontier = new PriorityQueue<Location>();
        frontier.Enqueue(start, 0);

        cameFrom[start] = start;
        costSoFar[start] = 0;

        while (frontier.Count > 0)
        {
            var current = frontier.Dequeue();

            if (current.Equals(goal))
            {
                break;
            }

            foreach (var next in graph.Neighbors(current))
            {
                double newCost = costSoFar[current]

```

```

        + graph.Cost(current, next);
    if (!costSoFar.ContainsKey(next)
        || newCost < costSoFar[next])
    {
        costSoFar[next] = newCost;
        double priority = newCost + Heuristic(next, goal);
        frontier.Enqueue(next, priority);
        cameFrom[next] = current;
    }
}
}
}

public class Test
{
    static void DrawGrid(SquareGrid grid, AStarSearch astar) {
        // Print out the cameFrom array
        for (var y = 0; y < 10; y++)
        {
            for (var x = 0; x < 10; x++)
            {
                Location id = new Location(x, y);
                Location ptr = id;
                if (!astar.cameFrom.TryGetValue(id, out ptr))
                {
                    ptr = id;
                }
                if (grid.walls.Contains(id)) { Console.Write("##"); }
                else if (ptr.x == x+1) { Console.Write("\u2192 "); }
                else if (ptr.x == x-1) { Console.Write("\u2190 "); }
                else if (ptr.y == y+1) { Console.Write("\u2193 "); }
                else if (ptr.y == y-1) { Console.Write("\u2191 "); }
                else { Console.Write("* "); }
            }
            Console.WriteLine();
        }
    }

    static void Main()
    {
        // Make "diagram 4" from main article
        var grid = new SquareGrid(10, 10);
        for (var x = 1; x < 4; x++)
        {
            for (var y = 7; y < 9; y++)
            {
                grid.walls.Add(new Location(x, y));
            }
        }
    }
}

```

```

    }
}
grid.forests = new HashSet<Location>
{
    new Location(3, 4), new Location(3, 5),
    new Location(4, 1), new Location(4, 2),
    new Location(4, 3), new Location(4, 4),
    new Location(4, 5), new Location(4, 6),
    new Location(4, 7), new Location(4, 8),
    new Location(5, 1), new Location(5, 2),
    new Location(5, 3), new Location(5, 4),
    new Location(5, 5), new Location(5, 6),
    new Location(5, 7), new Location(5, 8),
    new Location(6, 2), new Location(6, 3),
    new Location(6, 4), new Location(6, 5),
    new Location(6, 6), new Location(6, 7),
    new Location(7, 3), new Location(7, 4),
    new Location(7, 5)
};

// Run A*
var astar = new AStarSearch(grid, new Location(1, 4),
                             new Location(8, 5));

DrawGrid(grid, astar);
}
}

```

I haven't worked with C# much but the structure of the code is the same for my Python and C++ examples, and you can use that same structure in C#.

4 Algorithm changes

#

The version of Dijkstra's Algorithm and A* on my pages is slightly different from what you'll see in an algorithms or AI textbook.

The pure version of Dijkstra's Algorithm starts the priority queue with all nodes, and does not have early exit. It uses a "decrease-key" operation in the queue. It's fine in theory. But in practice...

1. By starting the priority with only the start node, we can keep it small, which makes it faster and use less memory.
2. With early exit, we almost never need to insert all the nodes into the queue, and we can return the path as soon as it's found.
3. By not putting all nodes into the queue at the start, most of the time we can use a cheap insert operation instead of the more expensive decrease-key operation.
4. By not putting all nodes into the queue at the start, we can handle situations where we do not even know all the nodes, or where the number of nodes is infinite.

This variant is sometimes called "Uniform Cost Search". See [Wikipedia](#)^[3] to see the pseudocode, or read [Felner's paper](#)^[4] [PDF] to see justifications for these changes.

There are three further differences between my version and what you might find elsewhere. These apply to both Dijkstra's Algorithm and A*:

5. I eliminate the check for a node being in the frontier with a higher cost. By not checking, I end up with duplicate elements in the frontier. *The algorithm still works.* It will revisit some locations more than necessary (but rarely, in my experience, as long as the heuristic is admissible). The code is simpler and it allows me to use a simpler and faster priority queue that does not support the decrease-key operation. The paper ["Priority Queues and Dijkstra's Algorithm"](#)^[5] suggests that this approach is faster in practice.
6. Instead of storing both a "closed set" and an "open set", I have call the open set the `frontier`, and I have a `reached` flag that tells me whether it's in *either* of those sets. I still have two sets but merging the two into `reached` simplifies the code.
7. I use hash tables instead of arrays of node objects. This eliminates the rather expensive *initialize* step that many other implementations have. For large game maps, the initialization of those arrays is often slower than the rest of A*.

If you have more suggestions for simplifications that preserve performance, please let me know!

5 Optimizations

#

For the code I present here, I've been focusing on simplicity and generality rather than performance. **First make it work, then make it fast.** Many of the optimizations I use in real projects are specific to the project, so instead of presenting optimal code, here are some ideas to pursue for your own project:

5.1 Graph

#

The biggest optimization you can make is to explore fewer nodes. My #1 recommendation is that if you're using a grid map, consider using a non-grid pathfinding graph. It's not always feasible but it's worth looking at.

If your graph has a simple structure (e.g. a grid), calculate the neighbors in a function. If it's a more complex structure (either a non-grid, or a grid with lots of walls, like a maze), store the neighbors in a data structure.

You can also save a bit of copying by reusing the neighbors array. Instead of *returning* a new one each time, allocate it once in the search code and pass it into the graph's neighbors method.

5.2 Queue

#

Breadth First Search uses a simple queue instead of the priority queue needed by the other algorithms. Queues are simpler and faster than priority queues. In exchange, the other algorithms usually explore fewer nodes. In most game maps, exploring fewer nodes is worth the slowdown from the other algorithms. There are some maps though where you don't save much, and it might be better to use Breadth First Search.

For queues, use a deque instead of an array. A deque allows fast insertion and removal on either end, whereas an array is fast only at one end. In Python, see [collections.deque](#)^[6]; in C++, see the [deque](#)^[7] container. However, breadth first search doesn't even need a queue; it can use two vectors, swapping them when one is empty.

For priority queues, use a binary heap instead of an array or sorted array. A binary heap allows fast insertion and removal, whereas an array is fast at one or the other but not both. In Python, see [heapq](#)^[8]; in C++, see the [priority_queue](#)^[9] container.

In Python, the Queue and PriorityQueue classes I presented above are so simple that you might consider inlining the methods into the search algorithm. I don't know if this buys you much; I need to measure it. The C++ versions are going to be inlined.

In Dijkstra's Algorithm, note that the priority queue's priority is stored twice, once in the priority queue and once in `cost_so_far`, so you could write a priority queue that gets priorities from elsewhere. I'm not sure if it's worth it.

The paper "[Priority Queues and Dijkstra's Algorithm](#)"^[10] by Chen, Chowdhury, Ramachandran, Lan Roche, Tong suggests optimizing the structure of Dijkstra's Algorithm by not reprioritizing, and it also suggests looking at [pairing heaps](#)^[11] and other data structures.

If you're considering using something other than a binary heap, first measure the size of your frontier and how often you reprioritize. Profile the code and see if the priority queue is the bottleneck.

My gut feeling is that *bucketing* is promising. Just as bucket sort and radix sort can be useful alternatives to quicksort when the keys are integers, we have an even better situation with Dijkstra's Algorithm and A*. The priorities in Dijkstra's Algorithm are *incredibly narrow*. If the lowest element in the queue has priority f , then the highest element has priority $f+e$ where e is the maximum edge weight. In the forest example, I have edge weights 1 and 5.

That means all the priorities in the queue are going to be between f and $f+5$. Since they're all integers, *there are only six different priorities*. We could use six buckets and not sort anything at all! A* produces a wider range of priorities but it's still worth looking at. And there are fancier bucketing approaches that handle a wider range of situations.

Note that if all the edge weights are 1, the priorities will all be between f and $f+1$. This yields a variant of Breadth First Search that uses two arrays instead of a queue, which I used [on my hex grid page](#)^[12]. If the weights are 1 or 2, you'll have three arrays; if the weights are 1, 2, or 3, you'll have four arrays; and so on.

[I have more note about priority queue data structures here](#)^[13].

5.3 Search

#

The heuristic adds complexity and cpu time. The goal though is to explore fewer nodes. In some maps (such as mazes), the heuristic may not add much information, and it may be better to use a simpler algorithm without a heuristic guide.

Some people use an *inadmissible* (overestimating) heuristic to speed up A* search. This seems reasonable. I haven't looked closely into its implications though. I believe (but don't know for sure) that some already-reached elements may need to be visited again even after they've been taken out of the frontier.

Some implementations *always* insert a new node into the open set, even if it's already there. You can avoid the potentially expensive step of checking whether the node is already in the open set. This will make your open set bigger/slower and you'll also end up evaluating more nodes than necessary. If the open-set test is expensive, it might still be worth it. However, in the code I've presented, I made the test cheap and I don't use this approach.

Some implementations *don't test* whether a new node is better than an existing node in the open set. This avoids a potentially expensive check. However, it also *can lead to a bug*. For some types of maps, you will not find the shortest path when you skip this test. In the code I've presented, I check this (`new_cost < cost_so_far`). The test is cheap because I made it cheap to look up `cost_so_far`.

5.4 Integer locations

#

If your graph uses integers as locations, consider using a simple array instead of a hash table for `cost_so_far`, `reached`, `came_from`, etc. Since `reached` is an array of booleans, you can use a bit vector. Initialize the `reached` bit vector for all ids, but leave `cost_so_far` and `came_from` uninitialized. Then only initialize on the first visit.

```
vector<uint16_t> reached(1 + maximum_node_id/16);

...

size_t index = node_id/16;
uint16_t bitmask = 1u << (node_id & 0xf);
if (!(reached[index] & bitmask)
    || new_cost < cost_so_far[next]) {
    reached[index] |= bitmask;
    ...
}
```

If you run only one search at a time, you can statically allocate and then reuse these arrays from one invocation to the next. Then keep an array of all indices that have been assigned to the bit vector, and then reset those on exit. For example:

```
static vector<uint16_t> reached(1 + maximum_node_id/16);
static vector<size_t> indices_to_clear;

...

size_t index = node_id/16;
uint16_t bitmask = 1u << (node_id & 0xf);
if (!(reached[index] & bitmask)
    || new_cost < cost_so_far[next]) {
```

```

        if (!reached[index]) {
            indices_to_clear.push_back(index);
        }
        reached[index] |= bitmask;
        ...
    }
    ...

    for (size_t index : indices_to_clear) {
        reached[index] = 0;
    }
    indices_to_clear.clear();

```

(Caveat: I haven't used or tested this code)

6 Troubleshooting

#

6.1 Wrong paths

#

If you're not getting a shortest path, try testing:

- Does your priority queue work correctly? Try stopping the search and dequeuing all the elements. They should all be in order.
- Does your heuristic ever overestimate the true distance? The `priority` of a new node should never be lower than the priority of its parent, unless you are overestimating the distance (you can do this but you won't get shortest paths anymore). Try setting the heuristic to 0. If a 0 heuristic fixes the paths, your heuristic is probably wrong. If a 0 heuristic doesn't help, then your graph search algorithm code probably has a bug.
- In a statically typed language, the cost, heuristic, and priority values need to have compatible types. The sample code on this page works with either integers or floating point types, but not all graphs and heuristics are limited to integer values. Since priorities are the sum of costs and heuristics, the priorities will need to be floating point if *either* costs or heuristics are floating point.

- The heuristic and costs need to have the same “units”. Try testing A* on a map with no walls. If the heuristic and movement costs match up, the priority should be the *same* along the entire path. If it isn’t, then your heuristic and costs probably don’t match up. When there are no obstacles and uniform movement costs, at each step, the heuristic should decrease by the same amount `cost_so_far` increases.

6.2 Ugly paths

#

The most common question I get when people run pathfinding on a grid is *why don't my paths look straight?* On a grid with uniform movement costs, there can be more than one shortest path of the same length. For example, in a 4-way movement grid, moving south 2 and east 2 could be any of these: `SSEE`, `SESE`, `SEES`, `ESSE`, `ESES`, `EESS`. The pathfinding algorithm is going to pick one, and it may not be the one you prefer. The path is *short* but it doesn’t *look* good. What can we do to favor good looking paths, like `SESE` or `ESES`?



Many ways to move south 2 east 2

- *Straighten* the paths using a “string pulling” algorithm: If the final path has points P, Q, R, S, and there’s a straight line from P to S, then follow that straight line instead of visiting Q and R.
- *Don’t use a grid*: tell A* only the places where you might turn, instead of every grid square; [read more here](#). Bonus: switching to a non-grid usually makes A* much faster.
- *Modify the A* algorithm* to support “any angle” paths: Theta*, Block A*, Field A*, or AnyA. See the paper [An Empirical Comparison of Any-Angle Path-Planning Algorithms](#)^[14] from Uras & Koenig.

- *Nudge* the paths when there's a tie towards better-looking paths, by adjusting the order of nodes in the queue. For 4-way movement I describe two hacks below. For 8-way movement, make sure your neighbors function returns the cardinal directions (N, E, S, W) earlier in the array than the diagonal directions (NW, SE, SW, NE).

The hacks don't work as well as the other three approaches but they're easy to implement, so I'll describe them here:

6.2.1 Checkerboard neighbor order

Breadth First Search is sensitive to the order in which it explores the neighbors of a tile. We normally go through the neighbors in a fixed order. Since Breadth First Search uses a first-in-first-out queue, it will pick the *first* path to a node.

If moving south 2 and east 2, there are many ways to get there: `SSEE`, `SESE`, `SEES`, `ESSE`, `ESES`, `EESS`. If East comes before South in the list of neighbors, then it will *always* explore east before it explores south, and end up choosing `EESS`. If South comes before East, then it will always explore south first, and end up choosing `SSEE`. We can see that problem in this larger example where the order is East, North, West, South:

```
from implementation import *
test with custom order([(+1, 0), (0, -1), (-1, 0), (0, +1)])
```

[illegible]

It moves east as far as possible before considering north or south.

```

from implementation import *
test_with_custom_order([(0, +1), (-1, 0), (0, -1), (+1, 0)])

```

```

from implementation import *
test_with_custom_order([(0, +1), (+1, 0), (-1, 0), (0, -1)])

```

Here's the hack for Breadth First Search: in the graph class, make the list of neighbors depend on $(x + y) \% 2$:

- when 0: return the South, North, West, East neighbors
- when 1: return the East, West, North, South neighbors

The result is the path alternating between vertical and horizontal steps:

```
from implementation import *
test_with_custom_order(None)
```

```
→ → → ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓ ← ← ← ← ← ##### . . . . .
→ → → → ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓ ← ← ← ← ← ##### . . . ↓ . .
→ → → → → ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓ ← ← ← ← ← ##### . . ↓ ↓ Z ↓
→ → ↑ ##### ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓ @ @ @ @ @ ← ← ← ← ← ##### . → → ↓ @ ↓
→ ↑ ↑ ##### ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓ @ @ ##### @ @ ← ← ← ← ← ##### → → → ↓ @ ↓
↑ ↑ ↑ ##### → ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓ @ @ ← ##### ↑ @ @ ← ← ##### ↓ @ ↓
↓ ↓ ↓ ##### → → ↓ @ @ ← ← ← ##### ↑ @ @ ← ← ##### ↓ @ ←
↓ ↓ ↓ ##### → → A @ ← ← ← ##### ↑ ↑ @ @ @ @ @ @ @ @ @ @ @ @ ←
↓ ↓ ↓ ##### → → ↑ ↑ ↑ ← ← ← ##### ↑ ↑ ↑ ↑ ↑ ← ← ← ← ← ← ← ← ←
↓ ↓ ↓ ##### → ↑ ↑ ↑ ↑ ↑ ↑ ↑ ↑ ↑ ##### ↑ ↑ ↑ ↑ ↑ ↑ ↑ ↑ ↑ ↑ ↑ ↑
→ ↓ ↓ ##### ↑ ↑ ↑ ↑ ↑ ↑ ↑ ↑ ↑ ##### ↑ ↑ ↑ ↑ ↑ ↑ ↑ ↑ ↑ ↑ ↑ ↑
→ → ↓ ##### ↑ ↑ ↑ ↑ ↑ ↑ ↑ ↑ ↑ ##### ↑ ↑ ↑ ↑ ↑ ↑ ↑ ↑ ↑ ↑ ↑ ↑
→ → → → → ↑ ↑ ↑ ↑ ↑ ↑ ↑ ↑ ↑ ##### ↑ ↑ ↑ ↑ ↑ ↑ ↑ ↑ ↑ ↑ ↑ ↑
→ → → → ↑ ↑ ↑ ↑ ↑ ↑ ↑ ↑ ↑ ##### ↑ ↑ ↑ ↑ ↑ ↑ ↑ ↑ ↑ ↑ ↑ ↑
→ → → ↑ ↑ ↑ ↑ ↑ ↑ ↑ ↑ ↑ ##### ↑ ↑ ↑ ↑ ↑ ↑ ↑ ↑ ↑ ↑ ↑ ↑
```

Here's the code:

```
class SquareGrid:
    ...
    def neighbors(self, id: GridLocation) -> Iterator[GridLocation]:
        (x, y) = id
        neighbors = [(x+1, y), (x-1, y), (x, y-1), (x, y+1)] # E W N S
        if (x + y) % 2 == 0: neighbors.reverse() # change to S N W E
        results = filter(self.in_bounds, neighbors)
        results = filter(self.passable, results)
        return results
```

This is a quick hack but it works with 4-way movement to make Breadth First Search paths look better. **I used this hack on these tutorial pages.** (Note: I came up with this hack for these tutorial pages; if you've seen a good reference please send it to me.)

6.2.2 Checkerboard movement costs

The direction order hack above works with Breadth First Search, but does it work with A*? Let's try:

```
from implementation import *

g = GridWithWeights(30, 15)
g.walls = DIAGRAM1_WALLS
start, goal = (8, 7), (27, 2)
came_from, cost_so_far = a_star_search(g, start, goal)
draw_grid(g, point_to=came_from, path=reconstruct_path(came_from, start, goal))
```

No, it doesn't. It changes the *insertion* order for the queue, but Dijkstra's Algorithm and A* use a *priority* queue that follows priority order instead of the insertion order. The priority in Dijkstra's Algorithm uses the movement cost; the priority in A* uses both the movement cost and the heuristic. We need to modify either the movement cost or the heuristic to change the priority order.

Here's the hack for A* and Dijkstra's Algorithm: in the graph class, make the movement cost depend on $(x + y) \% 2$:

- when 0: make horizontal movement slightly more expensive
- when 1: make vertical movement slightly more expensive

```
from implementation import *

g = GridWithAdjustedWeights(30, 15)
```

```

g.walls = DIAGRAM1_WALLS
start, goal = (8, 7), (27, 2)
came_from, cost_so_far = a_star_search(g, start, goal)
draw_grid(g, point_to=came_from, path=reconstruct_path(came_from, start, goal))

```

```

. . . . . ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓ ← ← ← ← ← ##### . . . . .
. . . . . ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓ ← ← ← ← ← ##### . . . . .
. . . . . → → ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓ ↓ ← ← ← ← ← ##### . . . . . Z
. . . ##### → ↓ ↓ ↓ ↓ ↓ ↓ ↓ @ @ @ @ @ @ @ ← ← ← ← ← ##### . . . ↓ @ ←
. . . ##### ↓ → ↓ ↓ ↓ ↓ ↓ @ @ ##### ↑ ← @ @ ↑ ← ##### . . . → ↓ @ ←
. . . ##### → ↓ → ↓ ↓ ↓ @ @ ← ##### ↑ ↑ ← @ @ ↑ ##### ↓ @ ←
. . . ##### → → ↓ ↓ @ @ ← ← ##### ↑ ↑ ↑ ← @ @ ##### ↓ @ ←
. . . ##### → → → A @ ← ← ← ##### ↑ ↑ ↑ ↑ ← @ @ @ @ @ @ @ @ ←
. . . ##### → → ↑ ↑ ↑ ← ↑ ← ##### ↑ ↑ ↑ ↑ ↑ ↑ ↑ ↑ ↑ ↑ ↑ ↑ .
. . . ##### → ↑ ↑ ↑ ↑ ↑ ↑ ← ↑ ##### . . . . . . . . . . .
. . . ##### . ↑ ↑ ↑ ↑ ↑ ↑ ↑ ← ##### . . . . . . . . . . .
. . . ##### . . ↑ ↑ ↑ ↑ ↑ ↑ ↑ ##### . . . . . . . . . . .
. . . . . . . . . . . . . . . . . . . . . . . . . . . . . . .
. . . . . . . . . . . . . . . . . . . . . . . . . . . . . . .
. . . . . . . . . . . . . . . . . . . . . . . . . . . . . . .

```

This works nicely with 4-way movement.

Here's the code:

```

class GridWithAdjustedWeights(GridWithWeights):
    def cost(self, from_node, to_node):
        prev_cost = super().cost(from_node, to_node)
        nudge = 0
        (x1, y1) = from_node
        (x2, y2) = to_node
        if (x1 + y1) % 2 == 0 and x2 != x1: nudge = 1
        if (x1 + y1) % 2 == 1 and y2 != y1: nudge = 1
        return prev_cost + 0.001 * nudge

```

This is a quick hack but it works with 4-way movement to make Dijkstra's Algorithm and A* paths look better. (Note: I came up with this hack for these tutorial pages; if you've seen this idea elsewhere please send me a reference so I can add it to the page.)

6.2.3 8-way movement

The above two hacks work for 4-way movement. What if you have 8-way movement? If all 8 directions have the same movement cost, we can end up with a path that takes diagonals when it seems like it shouldn't:

```
from implementation import *
test_with_custom_order([(-1, -1), (-1, +1), (+1, -1), (+1, +1), (+1, 0), (0
```

The 4-way tie-breaking hacks can be extended to work here:

- Breadth First Search: make sure the cardinal neighbors (N, S, E, W) come before the diagonal neighbors (NE, NW, SE, SW).
- Dijkstra's Algorithm, A*: add a tiny movement penalty (0.001) to diagonal movements.

After using one of these hacks, the path will look like this:

```
from implementation import *
test_with_custom_order([(+1, 0), (0, -1), (-1, 0), (0, +1), (-1, -1), (-1,
```

```

↑ ↑ ↑ ↑ ↑ ↑ ↑ ↑ ↑ ↑ ↑ ↑ ##### ↑ ↑ ↑ ↑ ↑ ↑ ↑ ↑ ↑ ↑ ↑ ↑
↑ ↑ ↑ ↑ ↑ ↑ ↑ ↑ ↑ ↑ ↑ ↑ ##### ↑ ↑ ↑ ↑ ↑ ↑ ↑ ↑ ↑ ↑ ↑ ↑

```

These hacks are easy to implement and give reasonable paths for grids. However, for even better paths, try the approaches listed in the [Ugly_paths](#) section.

7

Vocabulary

#

Algorithms textbooks often use mathematical notation with single-letter variable names. On these pages I've tried to use more descriptive variable names. Correspondences:

- `cost` is sometimes written as w or d or l or length
- `cost_so_far` is usually written as g or d or distance
- `heuristic` is usually written as h
- In A*, the `priority` is usually written as f , where $f = g + h$
- `came_from` is sometimes written as π or parent or previous or prev
- `frontier` is usually called OPEN or fringe
- the reached set is the union of OPEN and CLOSED
- locations such as `current` and `next` are called *states* or *nodes* and written with letters u, v

The OPEN, CLOSED, and reached sets are sets of states. These are not stored in their own data structures. Instead, they are contained as part of other data structures:

- The *elements* of the `frontier` are the OPEN set and the *priorities* of the `frontier` are the associated priority values.
- The *keys* of the `came_from` map are the reached set and the *values* of the `came_from` map are the parent pointers. Alternatively, if you want to keep the costs, the *keys* of the `cost_so_far` map are the reached set and the *values* of the `cost_so_far` map are the costs.
- `reached = OPEN \cup CLOSED`

We can reason about the OPEN, CLOSED, and reached sets even though they're not stored in a separate structure.

8 More reading

#

- Aleksander Nowak has written a **Go version** of this code at <https://github.com/vyrwu/a-star-redblob>^[15]
- Wikipedia links:
 - [Queue](#)^[16]
 - [Graph](#)^[17]
 - [Breadth-First Search](#)^[18]
 - (Greedy) [Best-First Search](#)^[19]
 - [Dijkstra's Algorithm](#)^[20]
 - [A* Algorithm](#)^[21]

Email me redblobgames@gmail.com, or tweet [@redblobgames](https://twitter.com/redblobgames), or comment:

Endnotes

[1]: <http://theory.stanford.edu/~amitp/GameProgramming/MovementCosts.html>

[2]: <http://theory.stanford.edu/~amitp/GameProgramming/MovementCosts.html>

[3]: https://en.wikipedia.org/wiki/Dijkstra%27s_algorithm#Practical_optimizations_and_infinite_graphs

[4]: <https://www.aaii.org/ocs/index.php/SOCS/SOCS11/paper/viewFile/4017/4357>

[5]: <http://www.cs.sunysb.edu/~rezaul/papers/TR-07-54.pdf>

[6]: <https://docs.python.org/3/library/collections.html>

[7]: <http://en.cppreference.com/w/cpp/container/deque>

[8]: <https://docs.python.org/2/library/heapq.html>

[9]: http://en.cppreference.com/w/cpp/container/priority_queue

- [10]: <http://www.cs.sunysb.edu/~rezaul/papers/TR-07-54.pdf>
- [11]: http://en.wikipedia.org/wiki/Pairing_heap
- [12]: <https://www.redblobgames.com/grids/hexagons/#range-obstacles>
- [13]: <http://theory.stanford.edu/~amitp/GameProgramming/ImplementationNotes.html#set-representation>
- [14]: https://scholar.google.com/scholar?cluster=8491292501067866547&hl=en&as_sdt=0,5
- [15]: <https://github.com/vyrwu/a-star-redblob>
- [16]: [http://en.wikipedia.org/wiki/Queue_\(abstract_data_type\)](http://en.wikipedia.org/wiki/Queue_(abstract_data_type))
- [17]: [http://en.wikipedia.org/wiki/Graph_\(data_structure\)](http://en.wikipedia.org/wiki/Graph_(data_structure))
- [18]: http://en.wikipedia.org/wiki/Breadth-first_search
- [19]: http://en.wikipedia.org/wiki/Best-first_search
- [20]: http://en.wikipedia.org/wiki/Dijkstra's_algorithm
- [21]: http://en.wikipedia.org/wiki/A*_search_algorithm

Copyright © 2020 *[Red Blob Games](#)*

 [RSS Feed](#)

Created with [Emacs Org-mode](#), from implementation.org. Last modified:
15 Oct 2020